



STORM @TWITTER

KARTHIK RAMASAMY

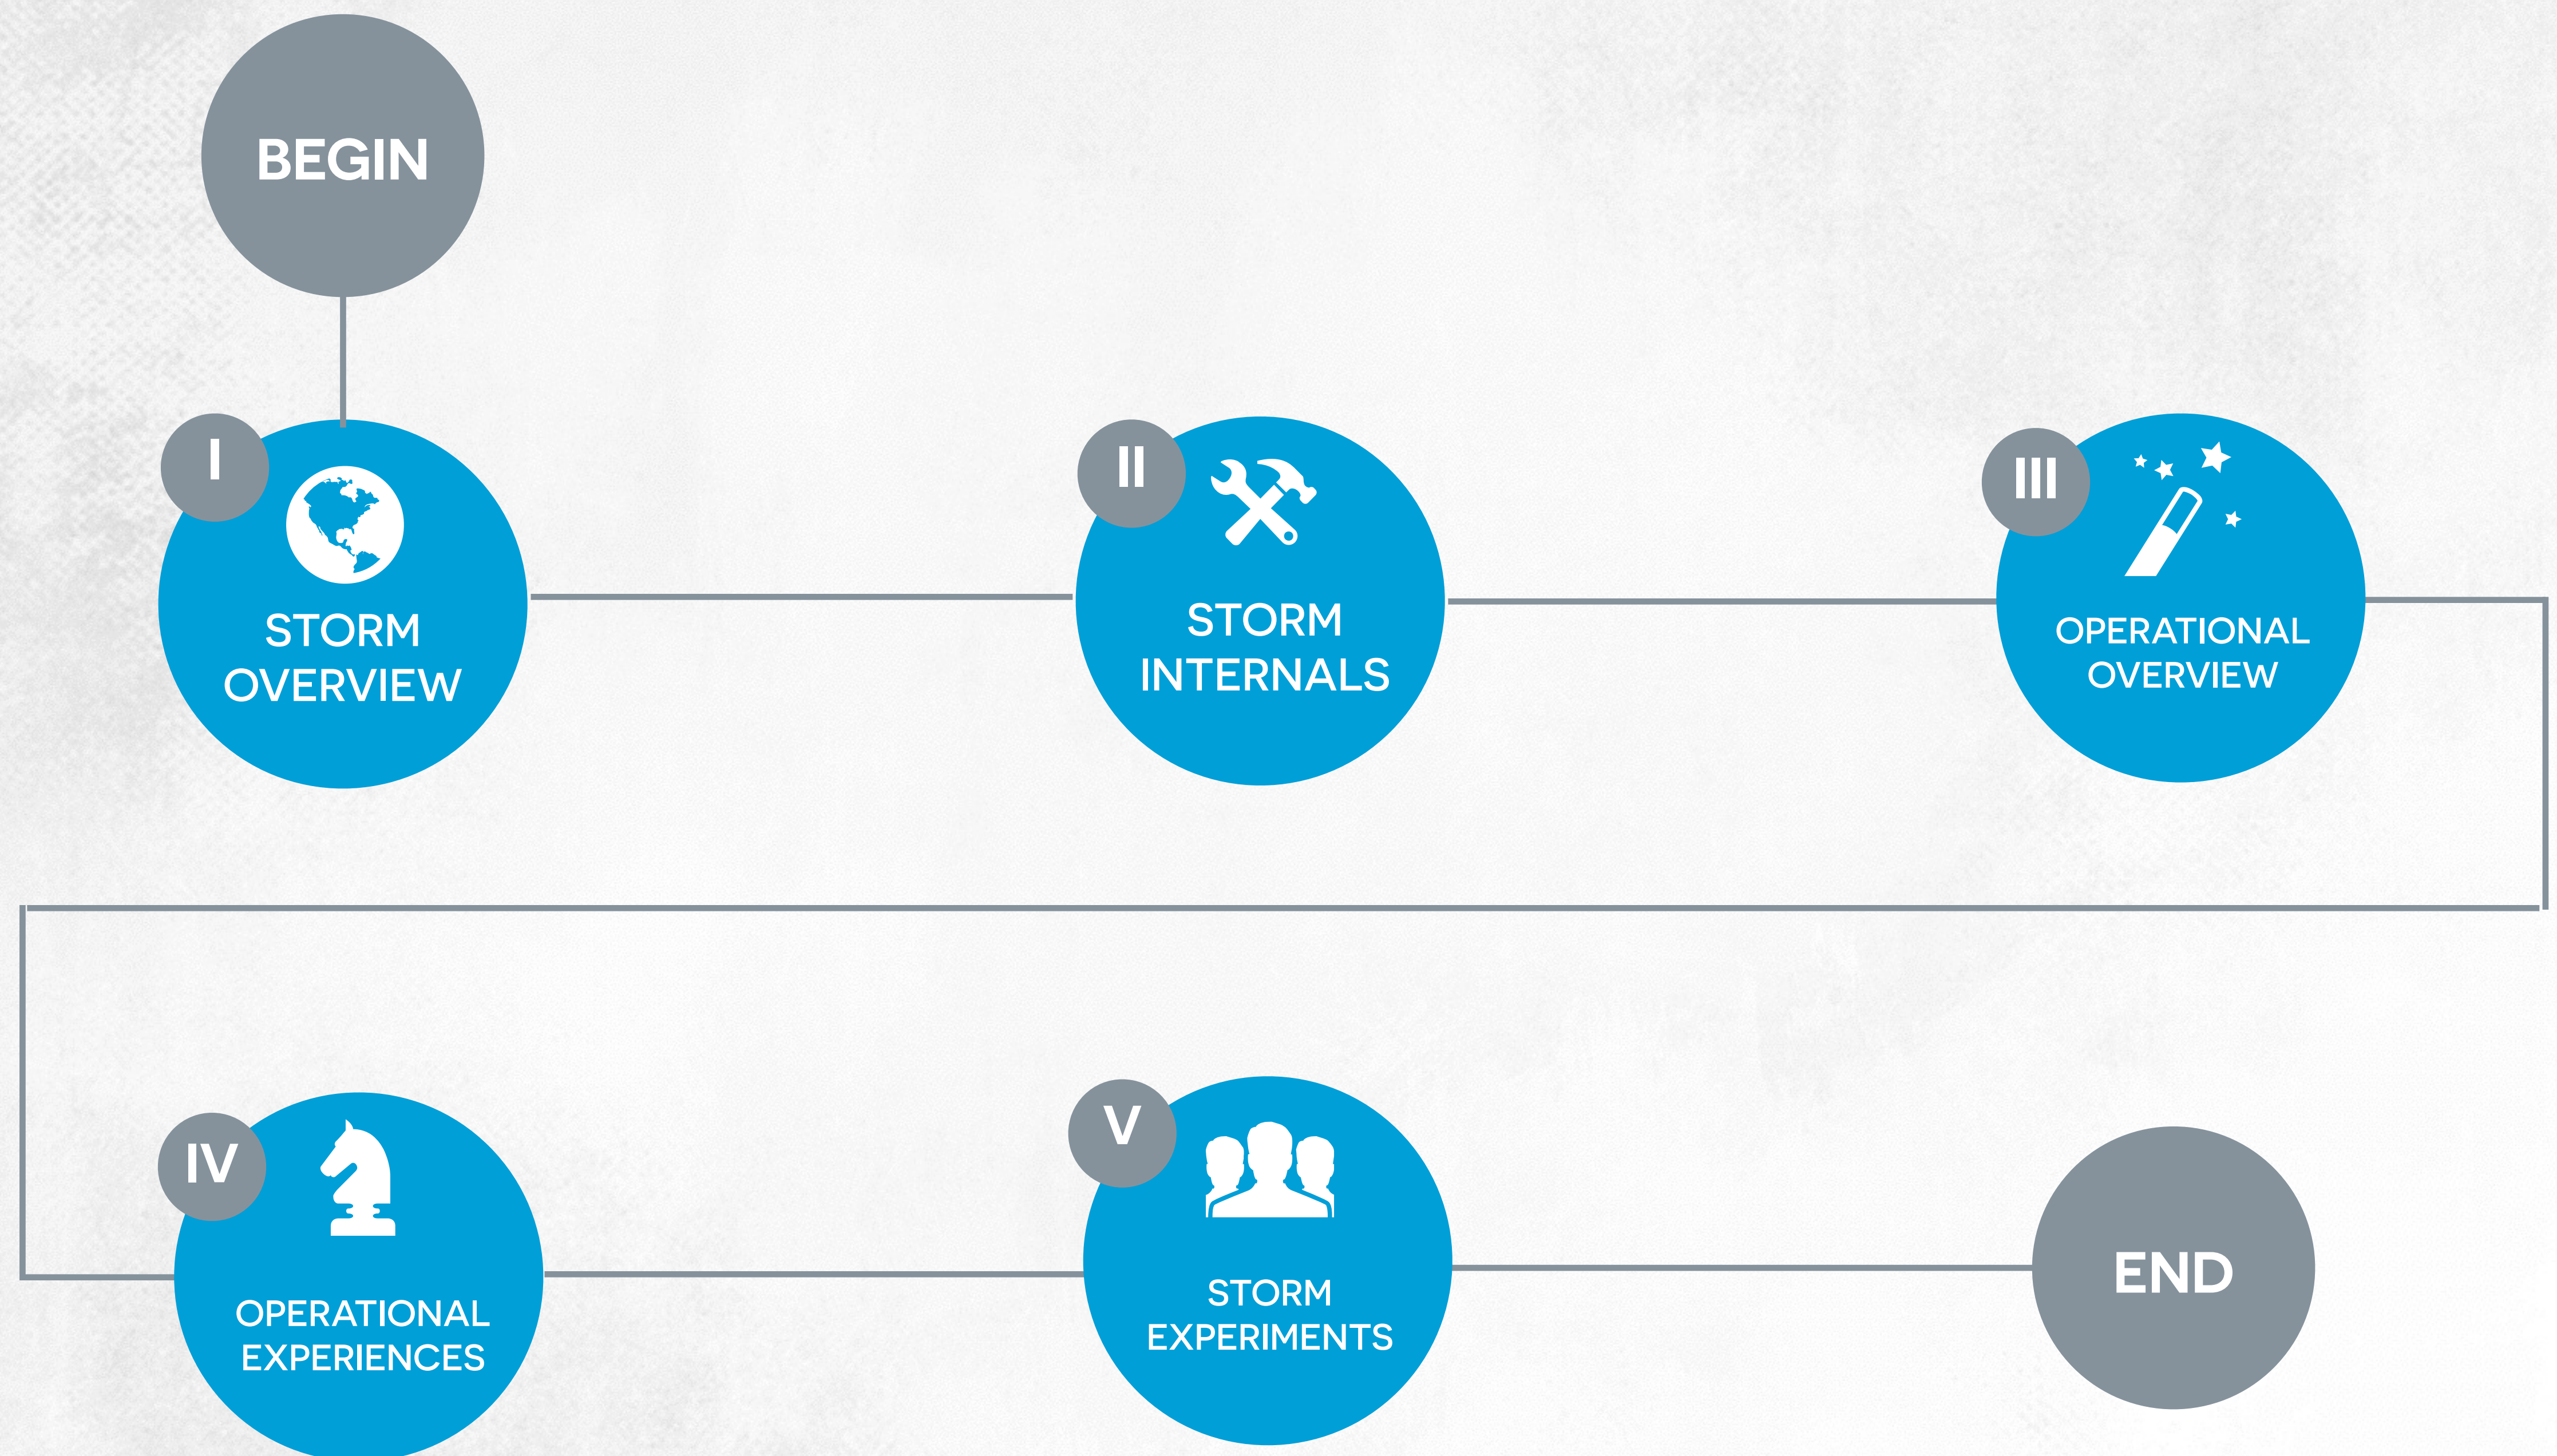
@KARTHIKZ

#TwitterAtSigmod

#TwitterDataStorm

Ankit Toshniwal, Siddarth Taneja, Amit Shukla, Jignesh Patel, Sanjeev Kulkarni
Jason Jackson, Krishna Gade, Maosong Fu, Jake Donham, Nikunj Bhagat
Sailesh Mittal and Dmitriy Ryaboy

TALK OUTLINE





OVERVIEW




WHAT IS **STORM**?


Streaming platform for analyzing realtime data as they arrive,
so you can react to data as it happens.




GUARANTEED
MESSAGE
PROCESSING



HORIZONTAL
SCALABILITY



ROBUST
FAULT
TOLERANCE



CONCISE
CODE- FOCUS
ON LOGIC



STORM DATA MODEL



TOPOLOGY

Directed acyclic graph

Vertices=computation, and edges=streams of data tuples



SPOUTS

Sources of data tuples for the topology

Examples – Kafka/Kestrel/MySQL/Postgres



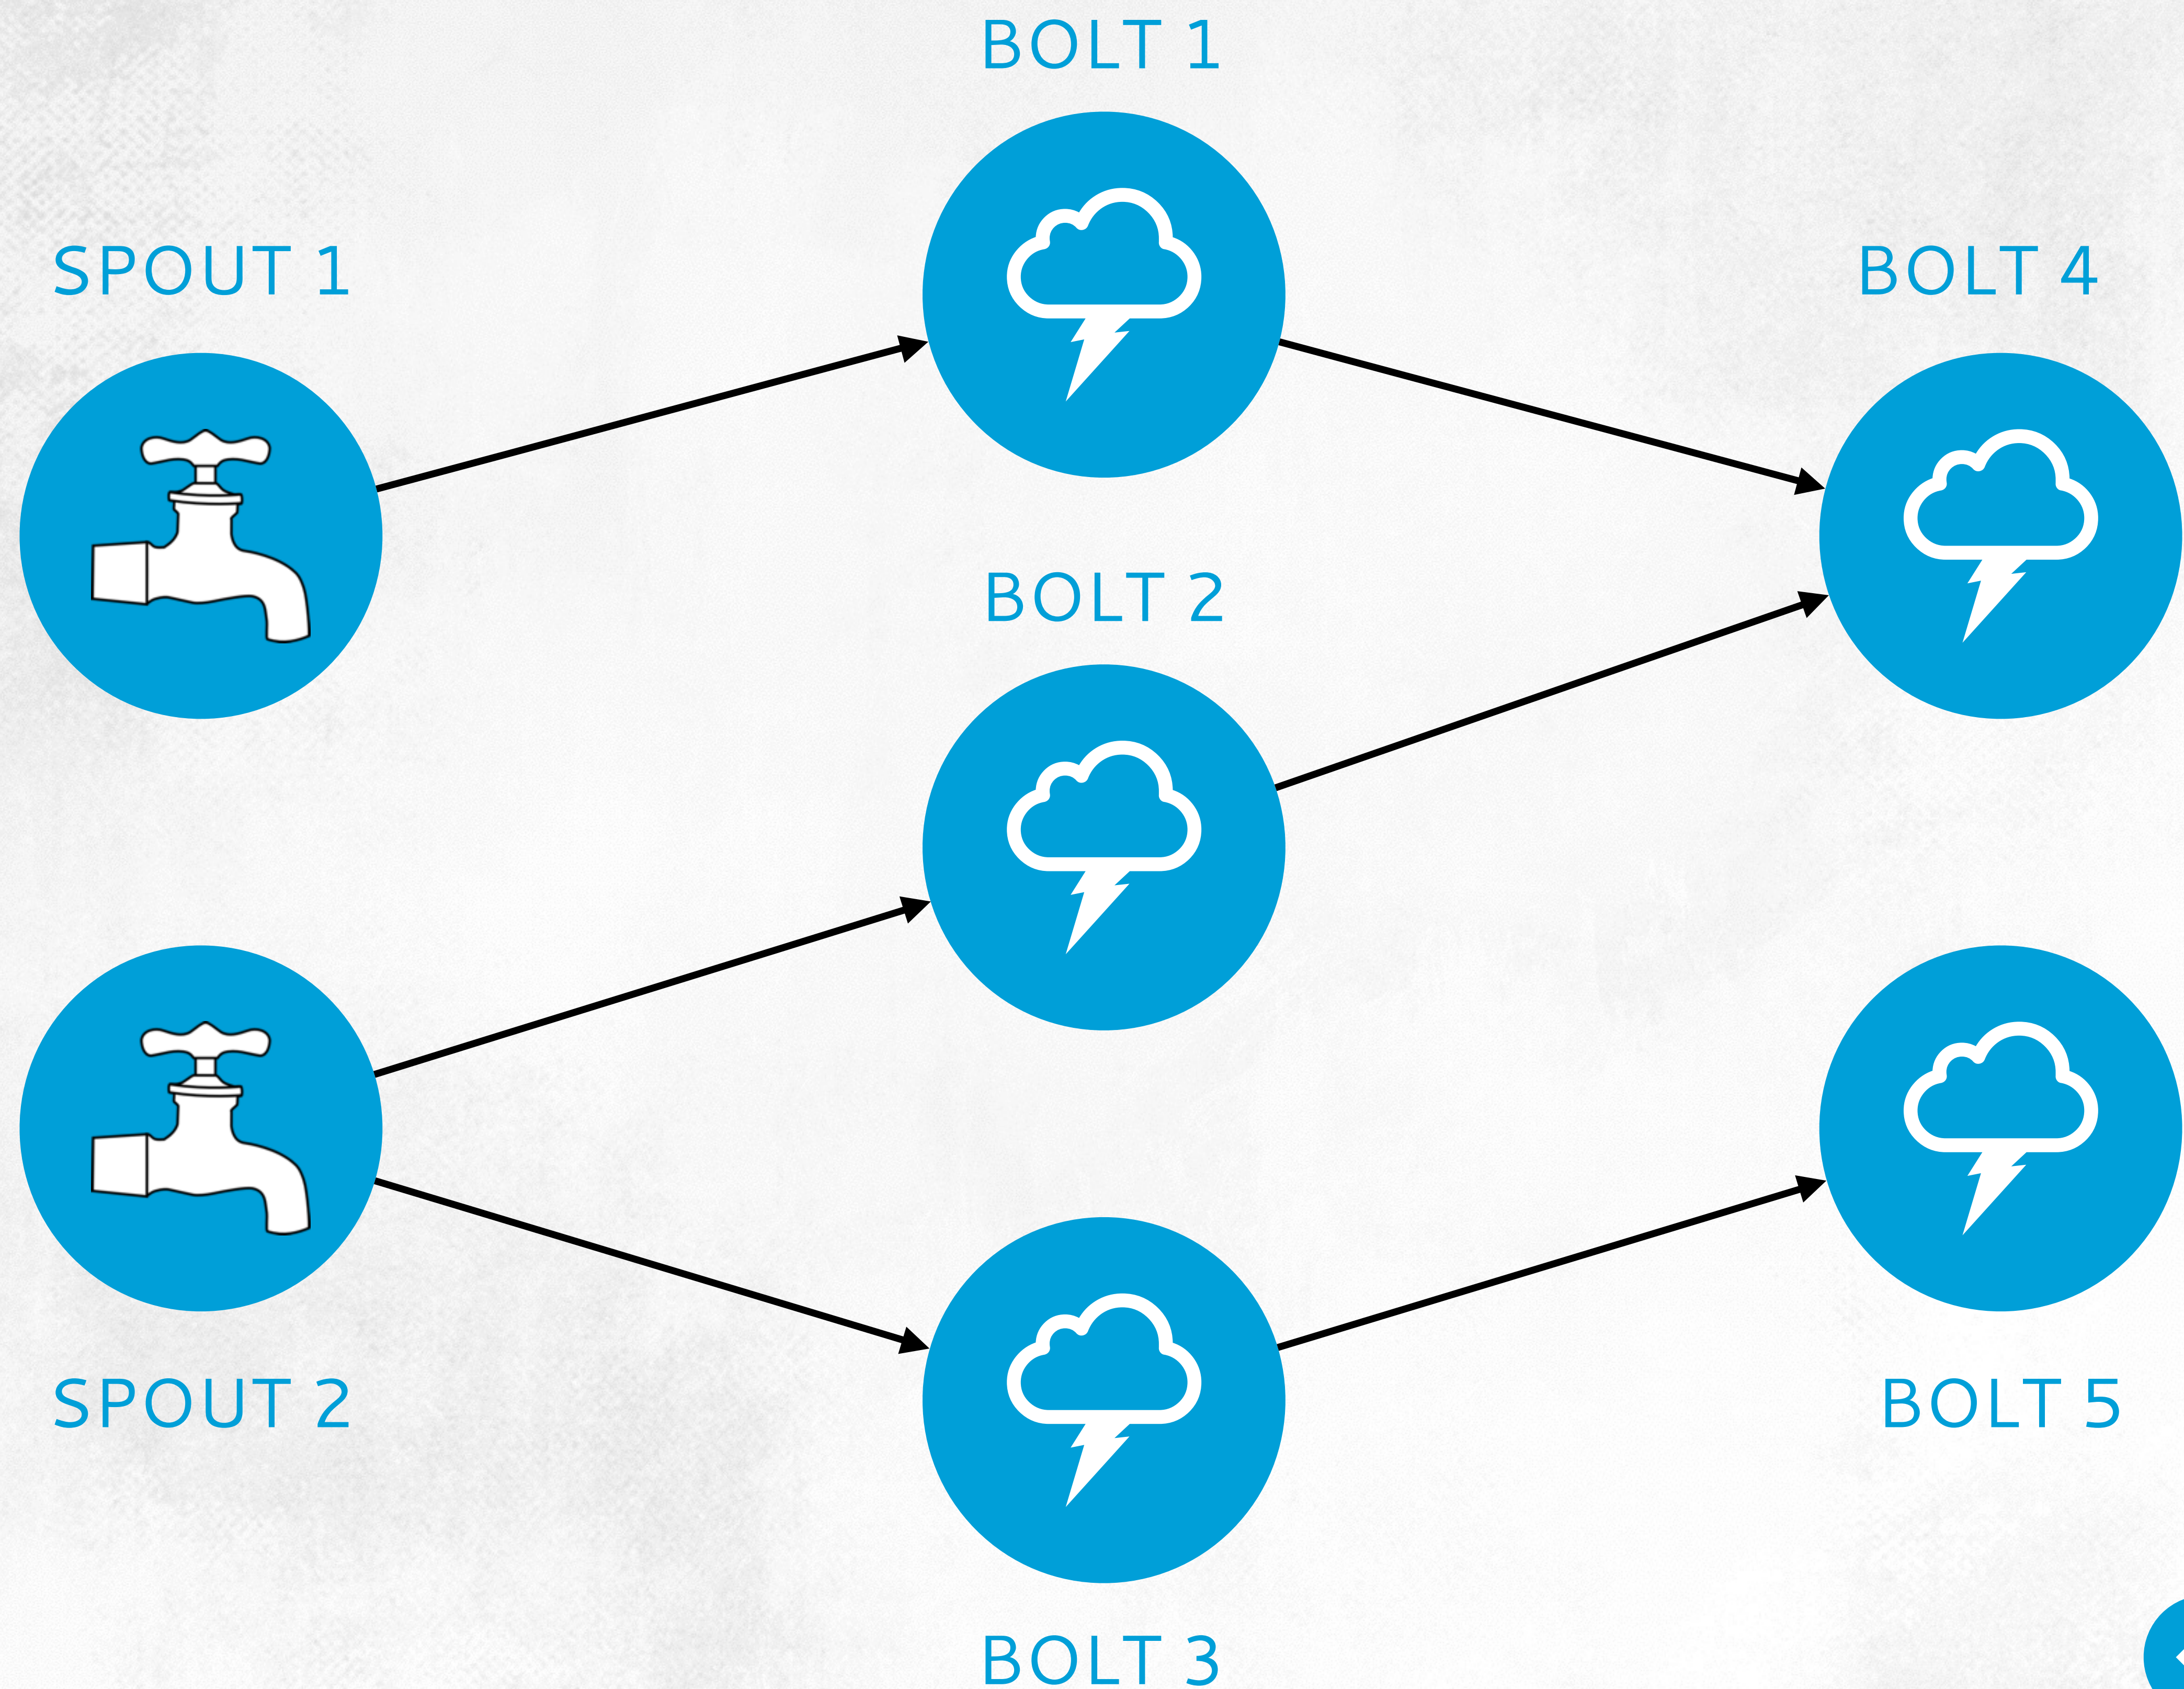
BOLTS

Process incoming tuples and emit outgoing tuples

Examples – filtering/aggregation/join/arbitrary function



STORM TOPOLOGY




WORD COUNT TOPOLOGY

Live stream of Tweets


#worldcup : 1M
soccer: 400K
....




TWEET SPOUT




PARSE TWEET BOLT



WORD COUNT BOLT



WORD COUNT TOPOLOGY




When a parse tweet bolt task emits a tuple
which word count bolt task should it send to?




STREAM GROUPINGS

SHUFFLE GROUPING




Random distribution
of tuples

FIELDS GROUPING




Group tuples by a
field or multiple
fields

ALL GROUPING



Replicates tuples to
all tasks

GLOBAL GROUPING



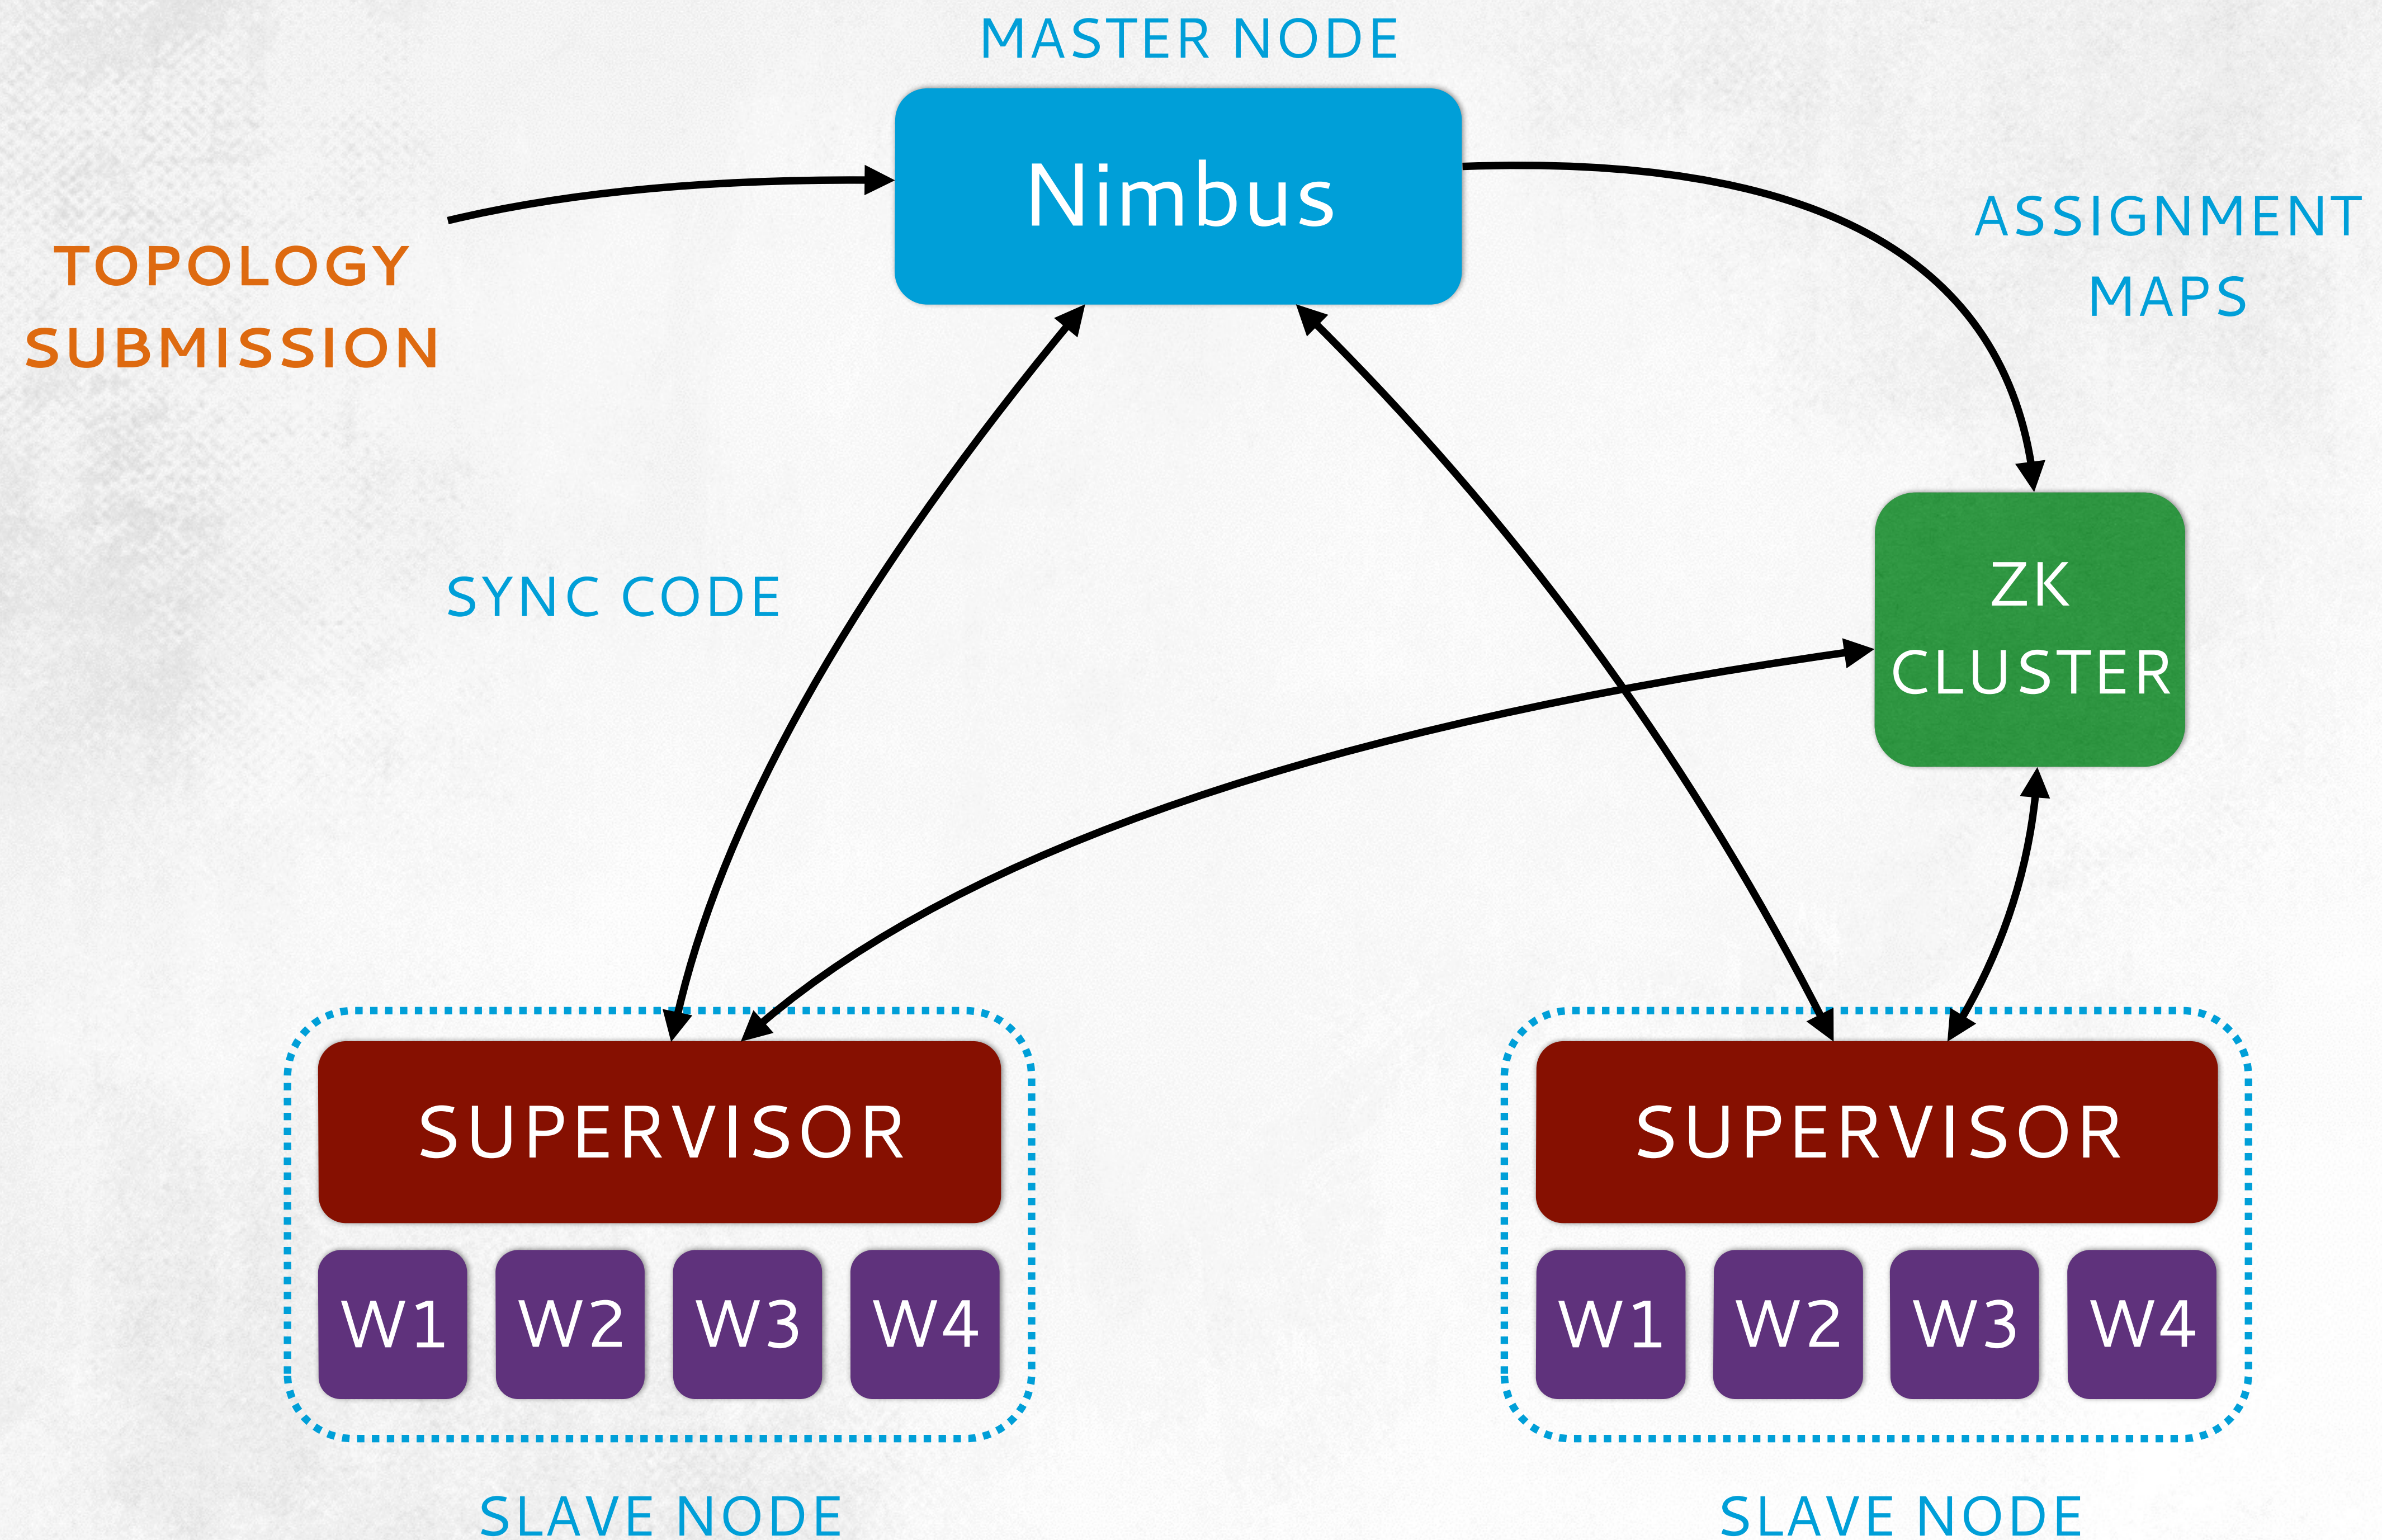
Sends the entire
stream to one task





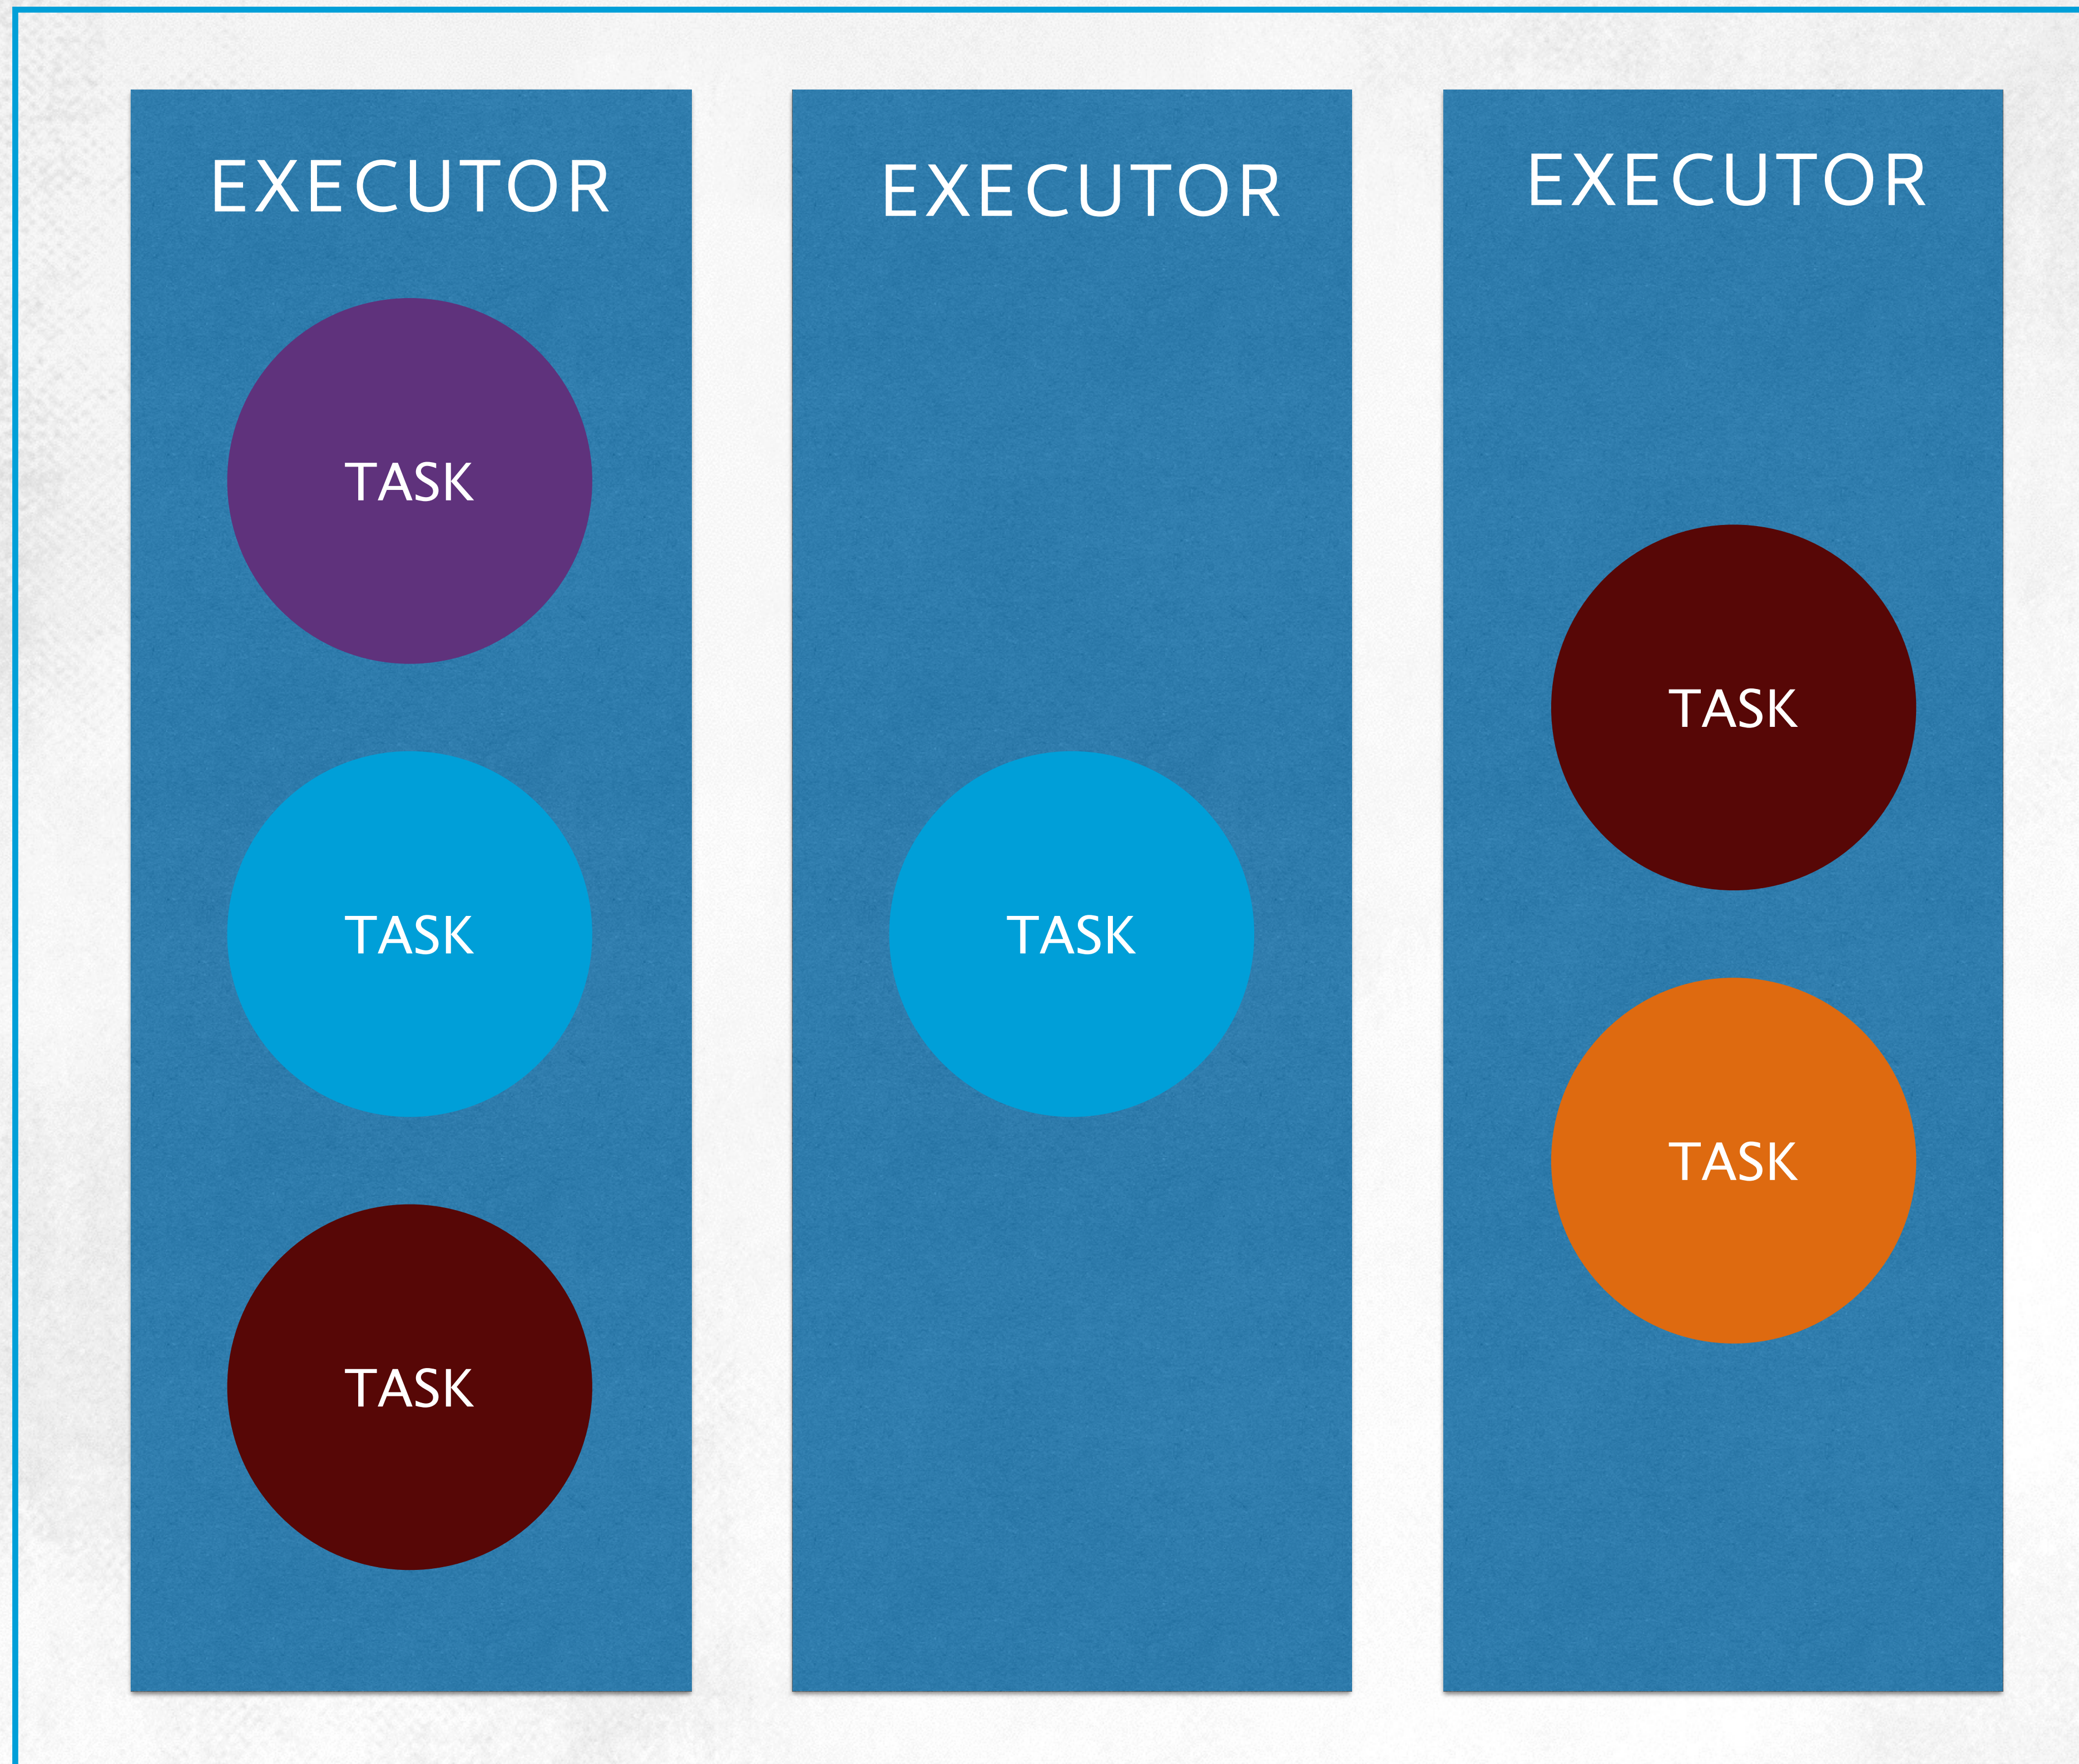
STORM INTERNALS

STORM ARCHITECTURE




STORM WORKER

JVM PROCESS



DATA FLOW IN STORM WORKERS





OPERATIONAL OVERVIEW

STORM METRICS

1

SUPPORT AND TROUBLE SHOOTING

2

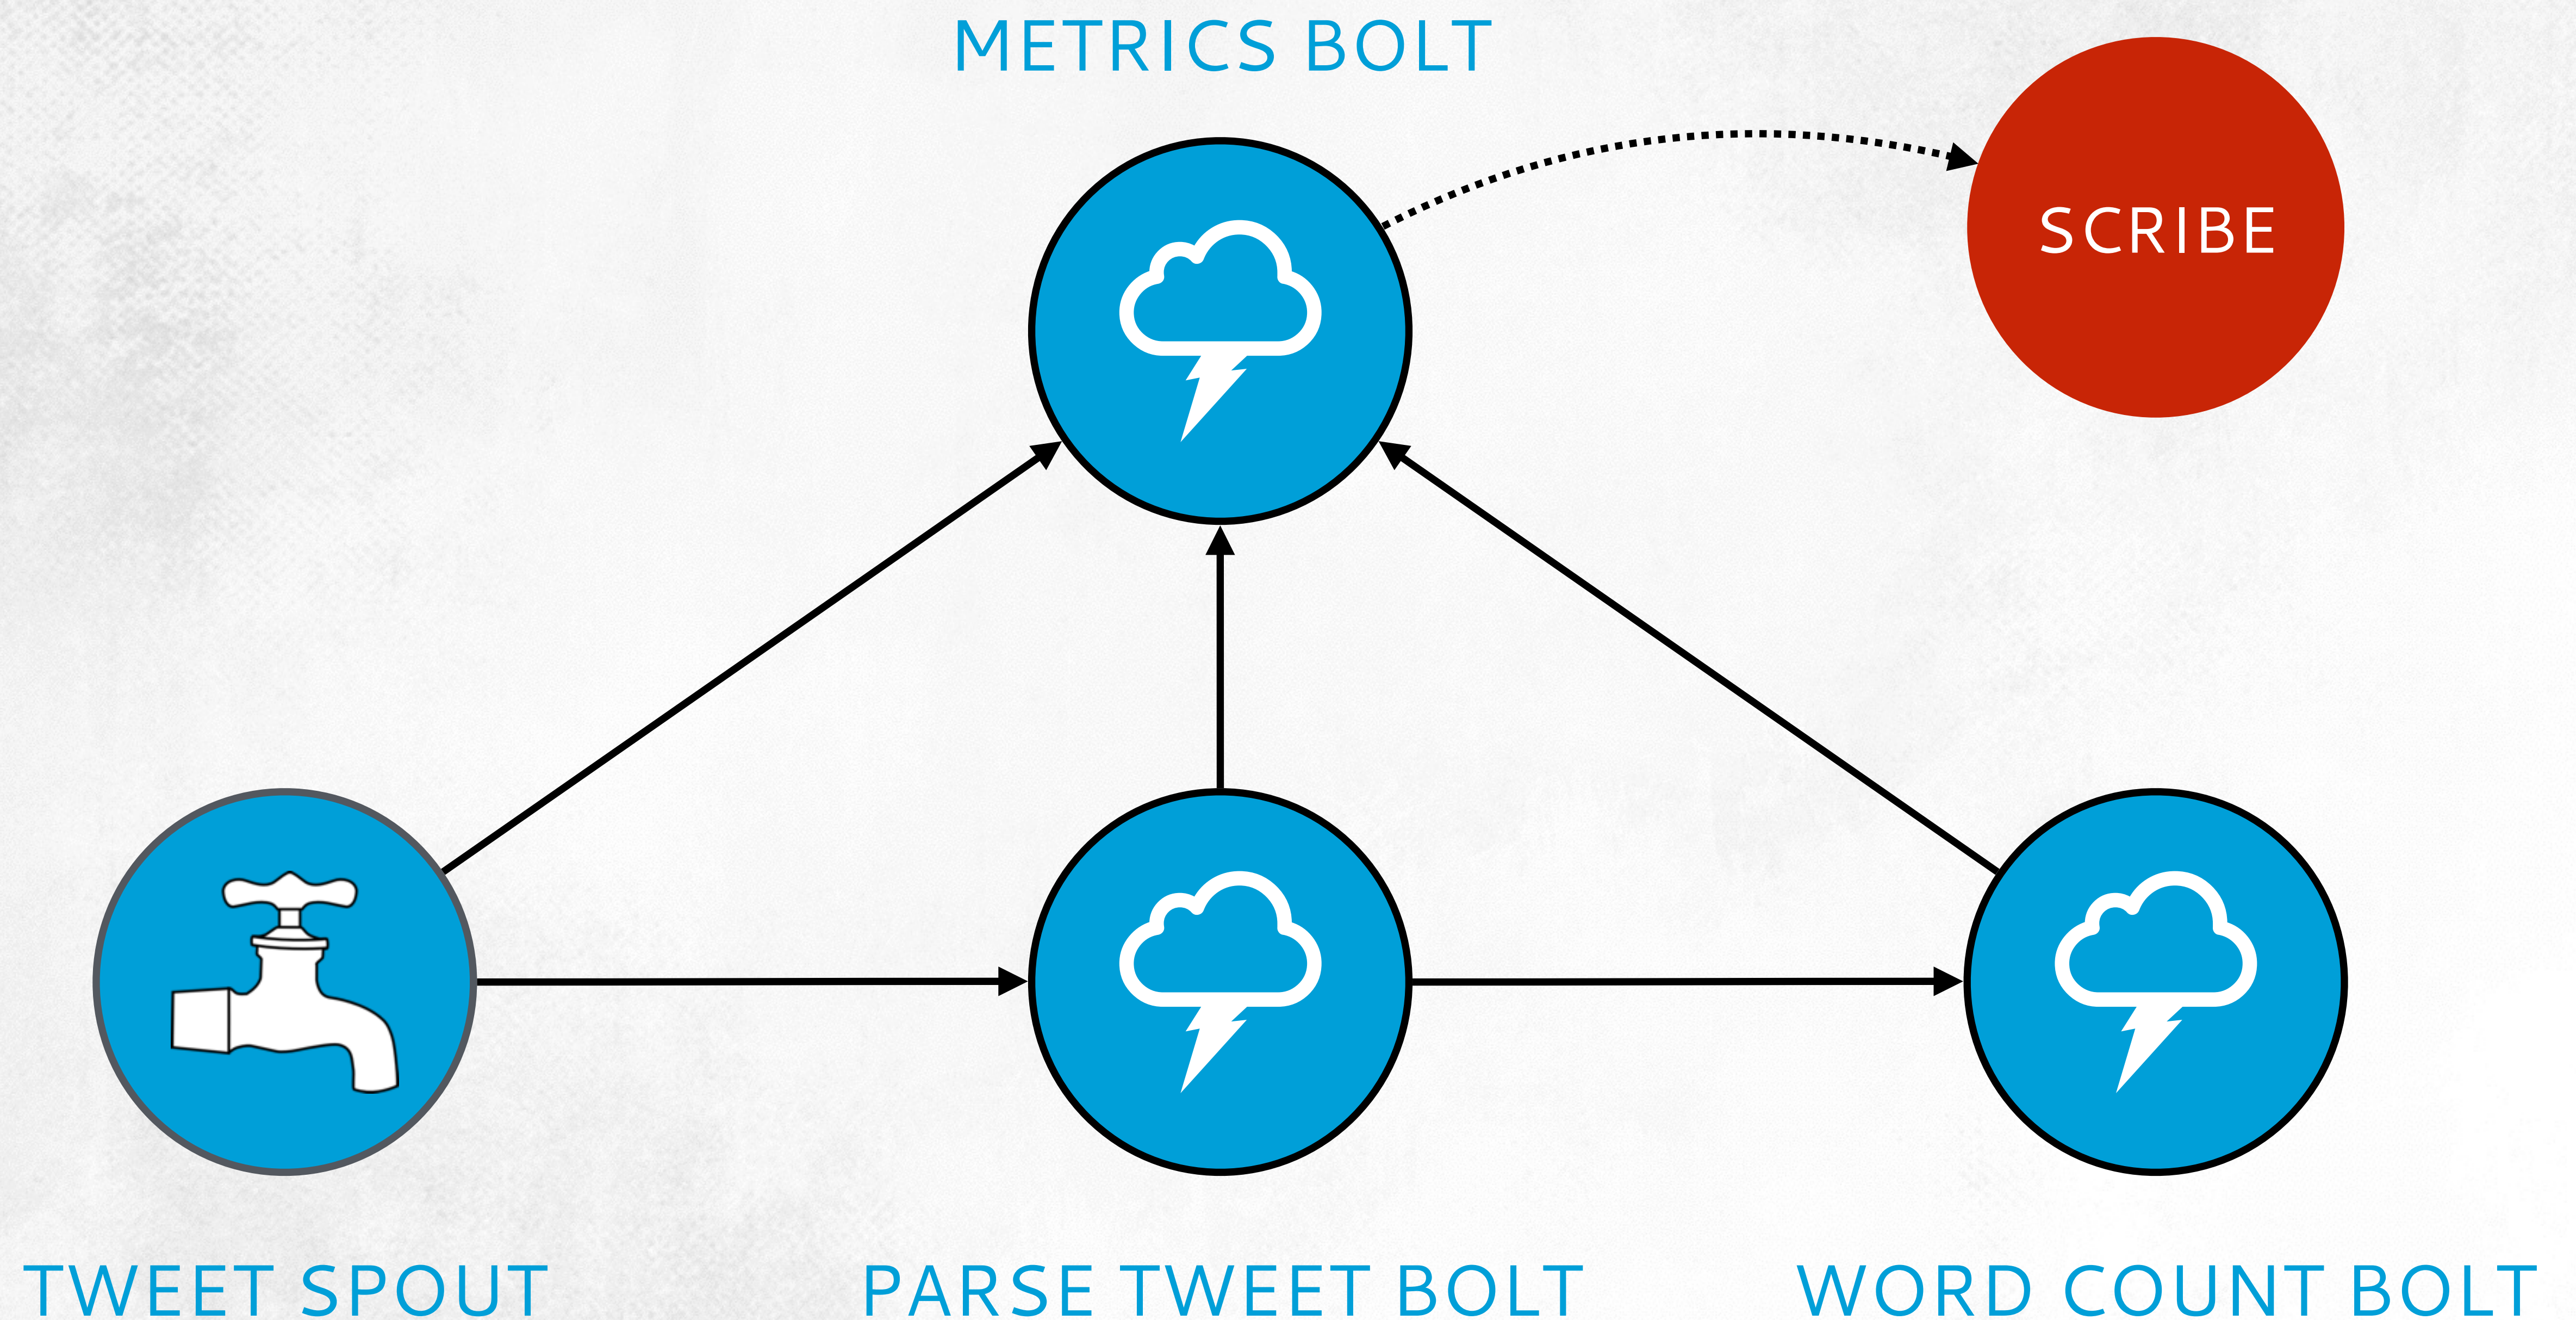
CONTINUOUS PERFORMANCE

3

CLUSTER AVAILABILITY



COLLECTING TOPOLOGY METRICS



SAMPLE TOPOLOGY DASHBOARD

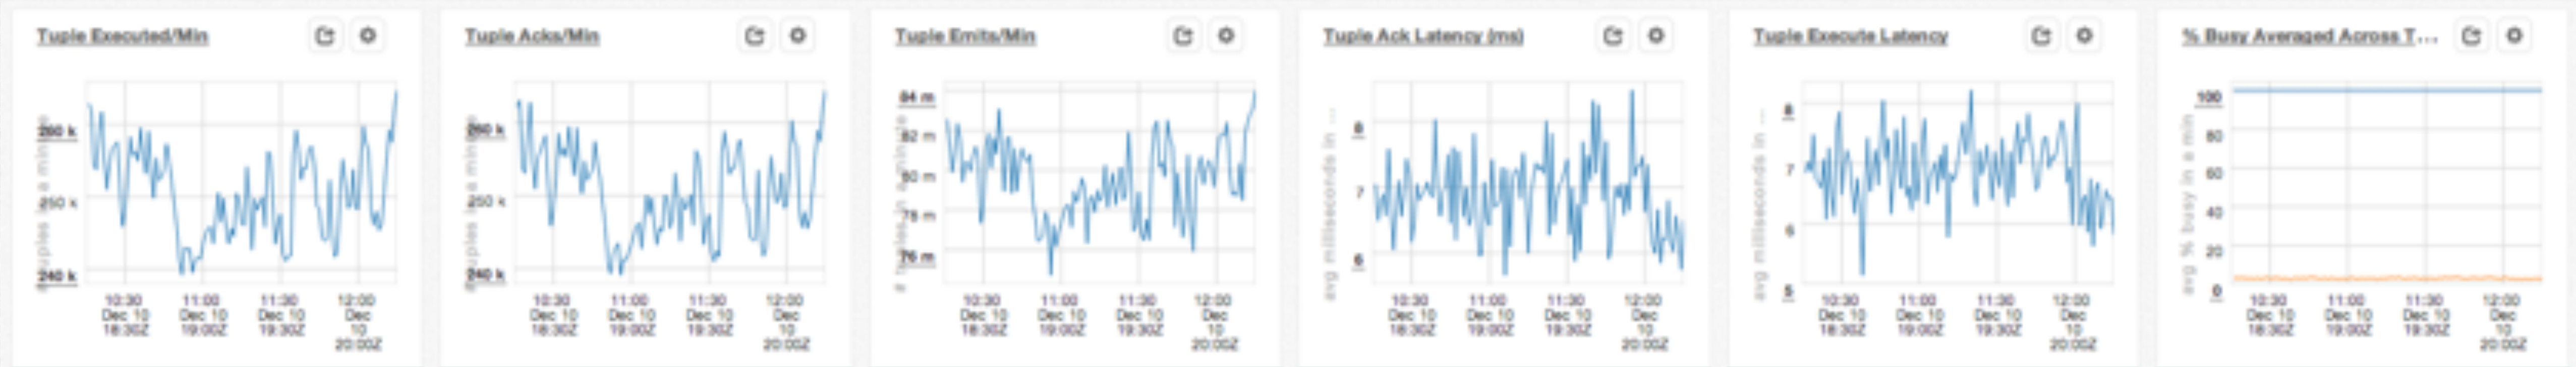
Workers (aka JVM Processes)



Spout: Tail-FlatMap-Source



Bolt: Tail-FlatMap





IV

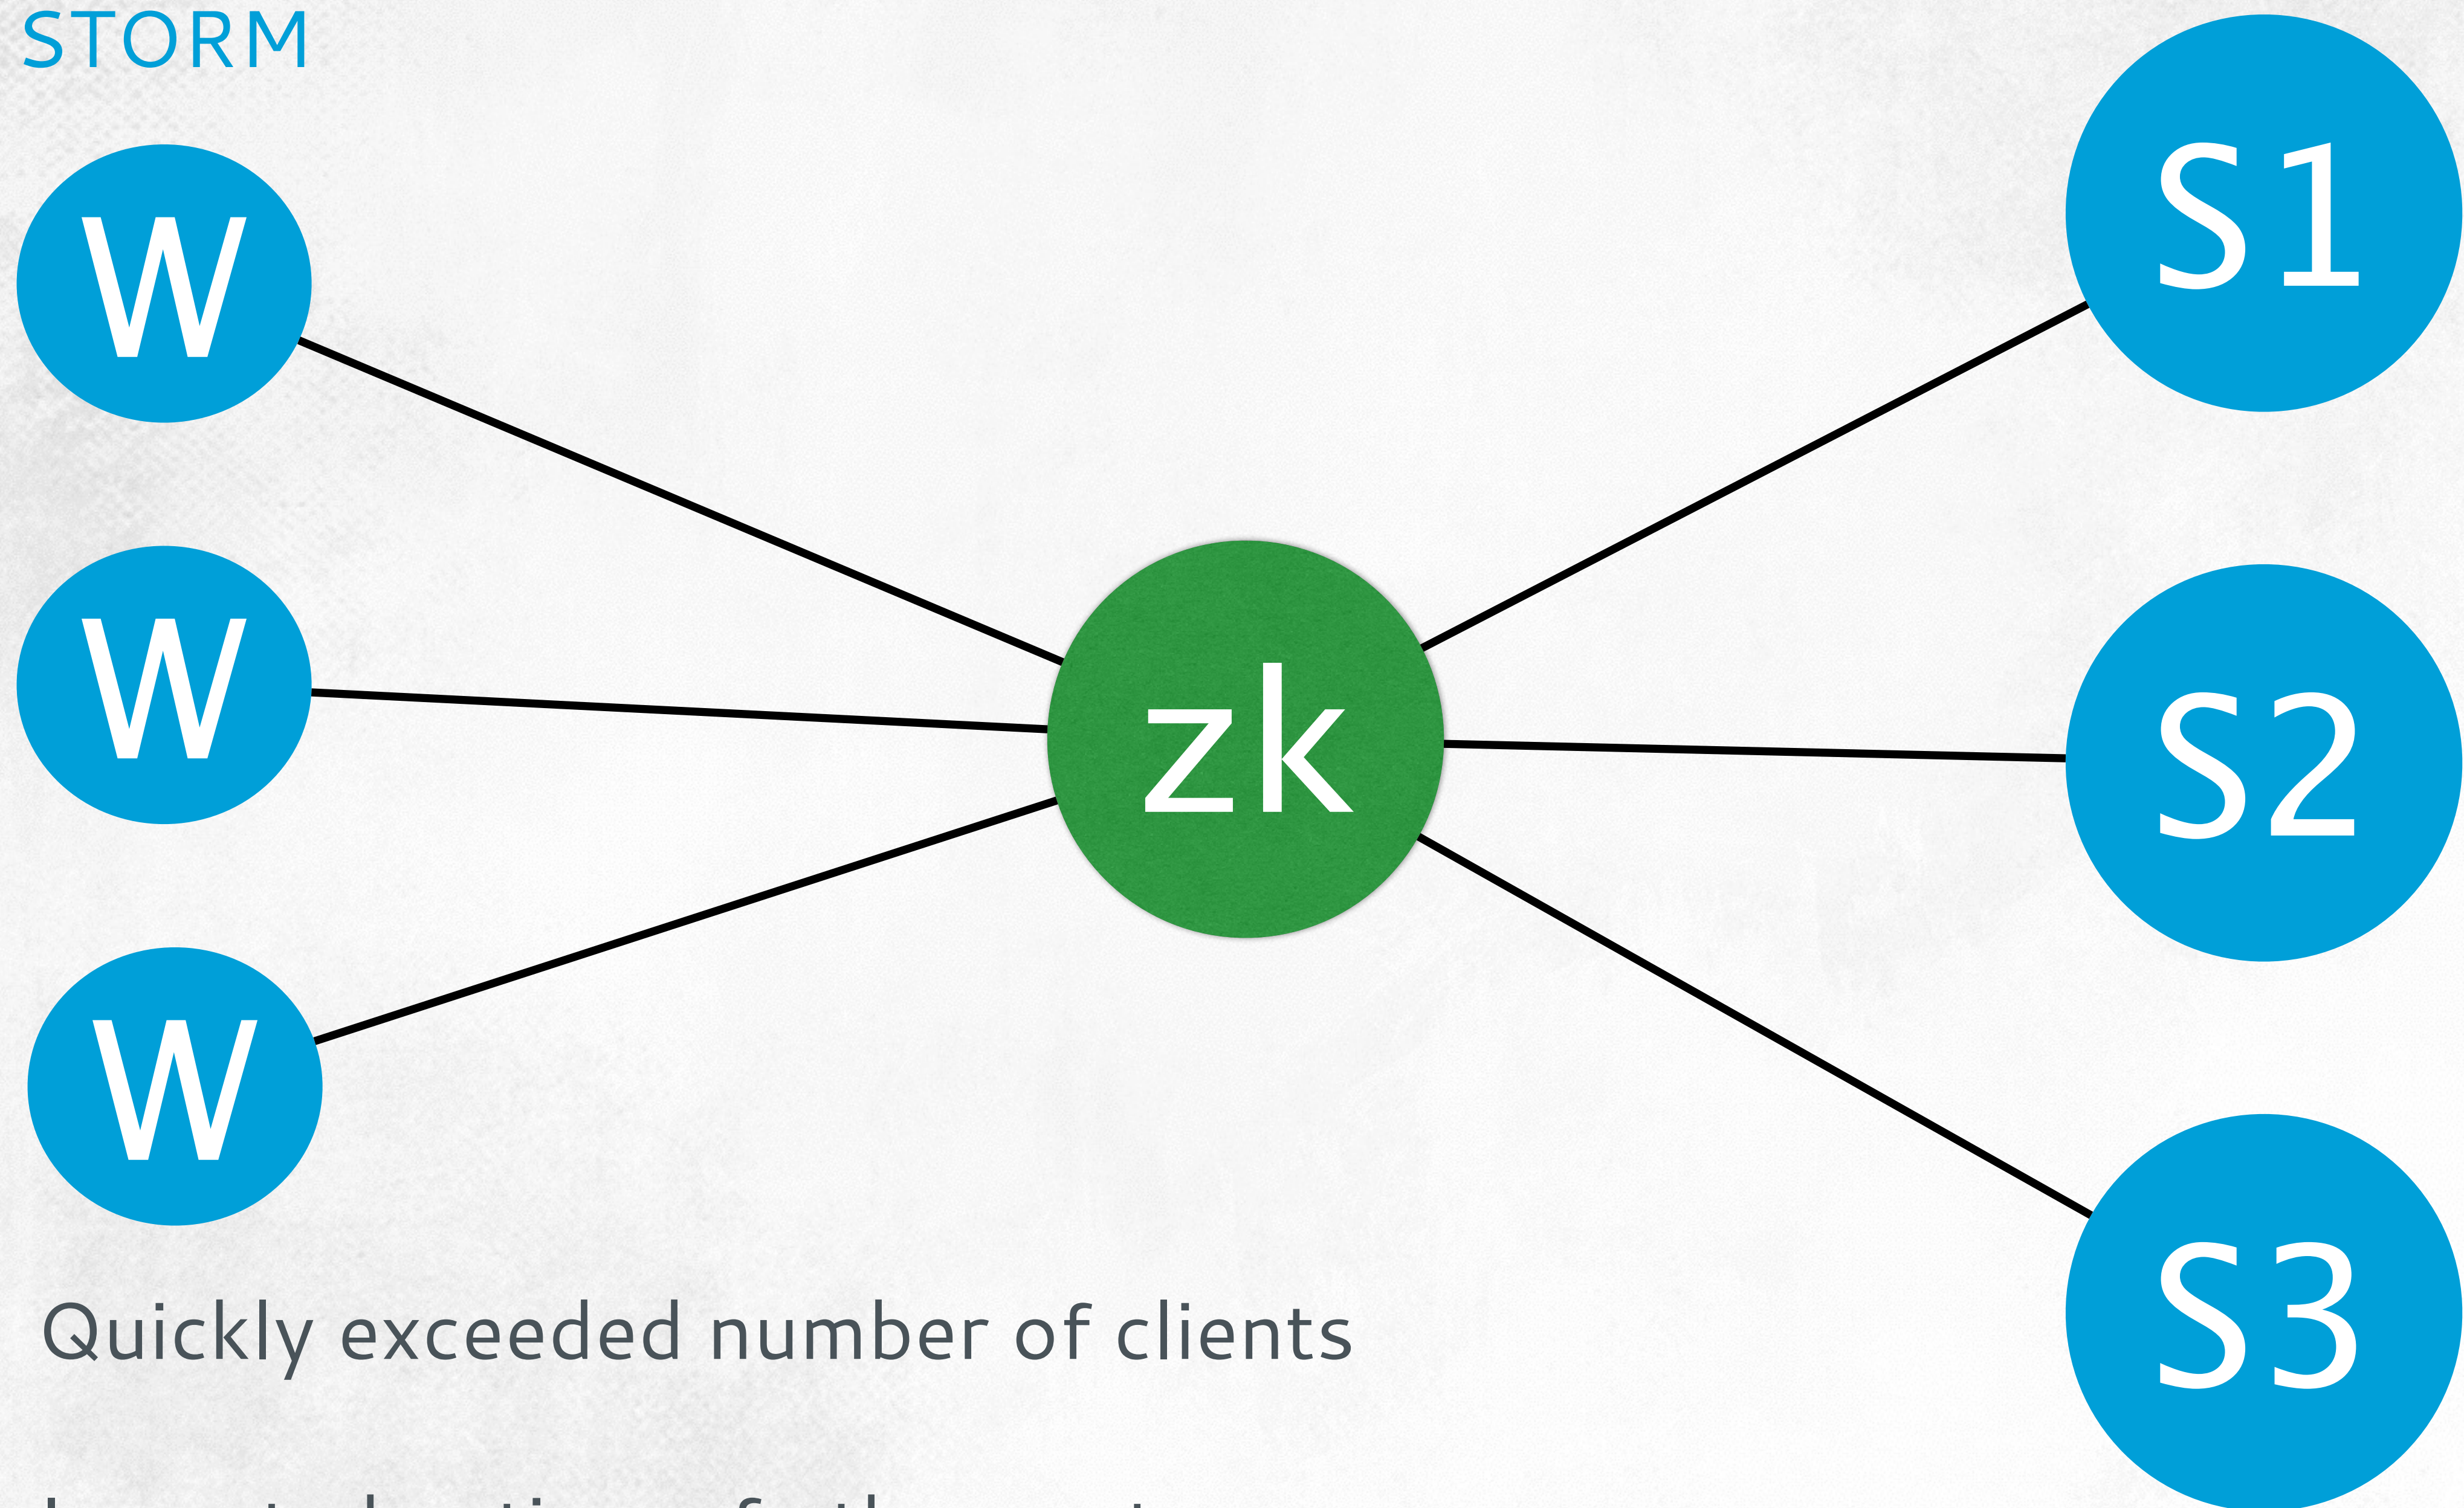
OPERATIONAL EXPERIENCES



OVERLOADED ZOOKEEPER

Shared configuration

STORM



Quickly exceeded number of clients

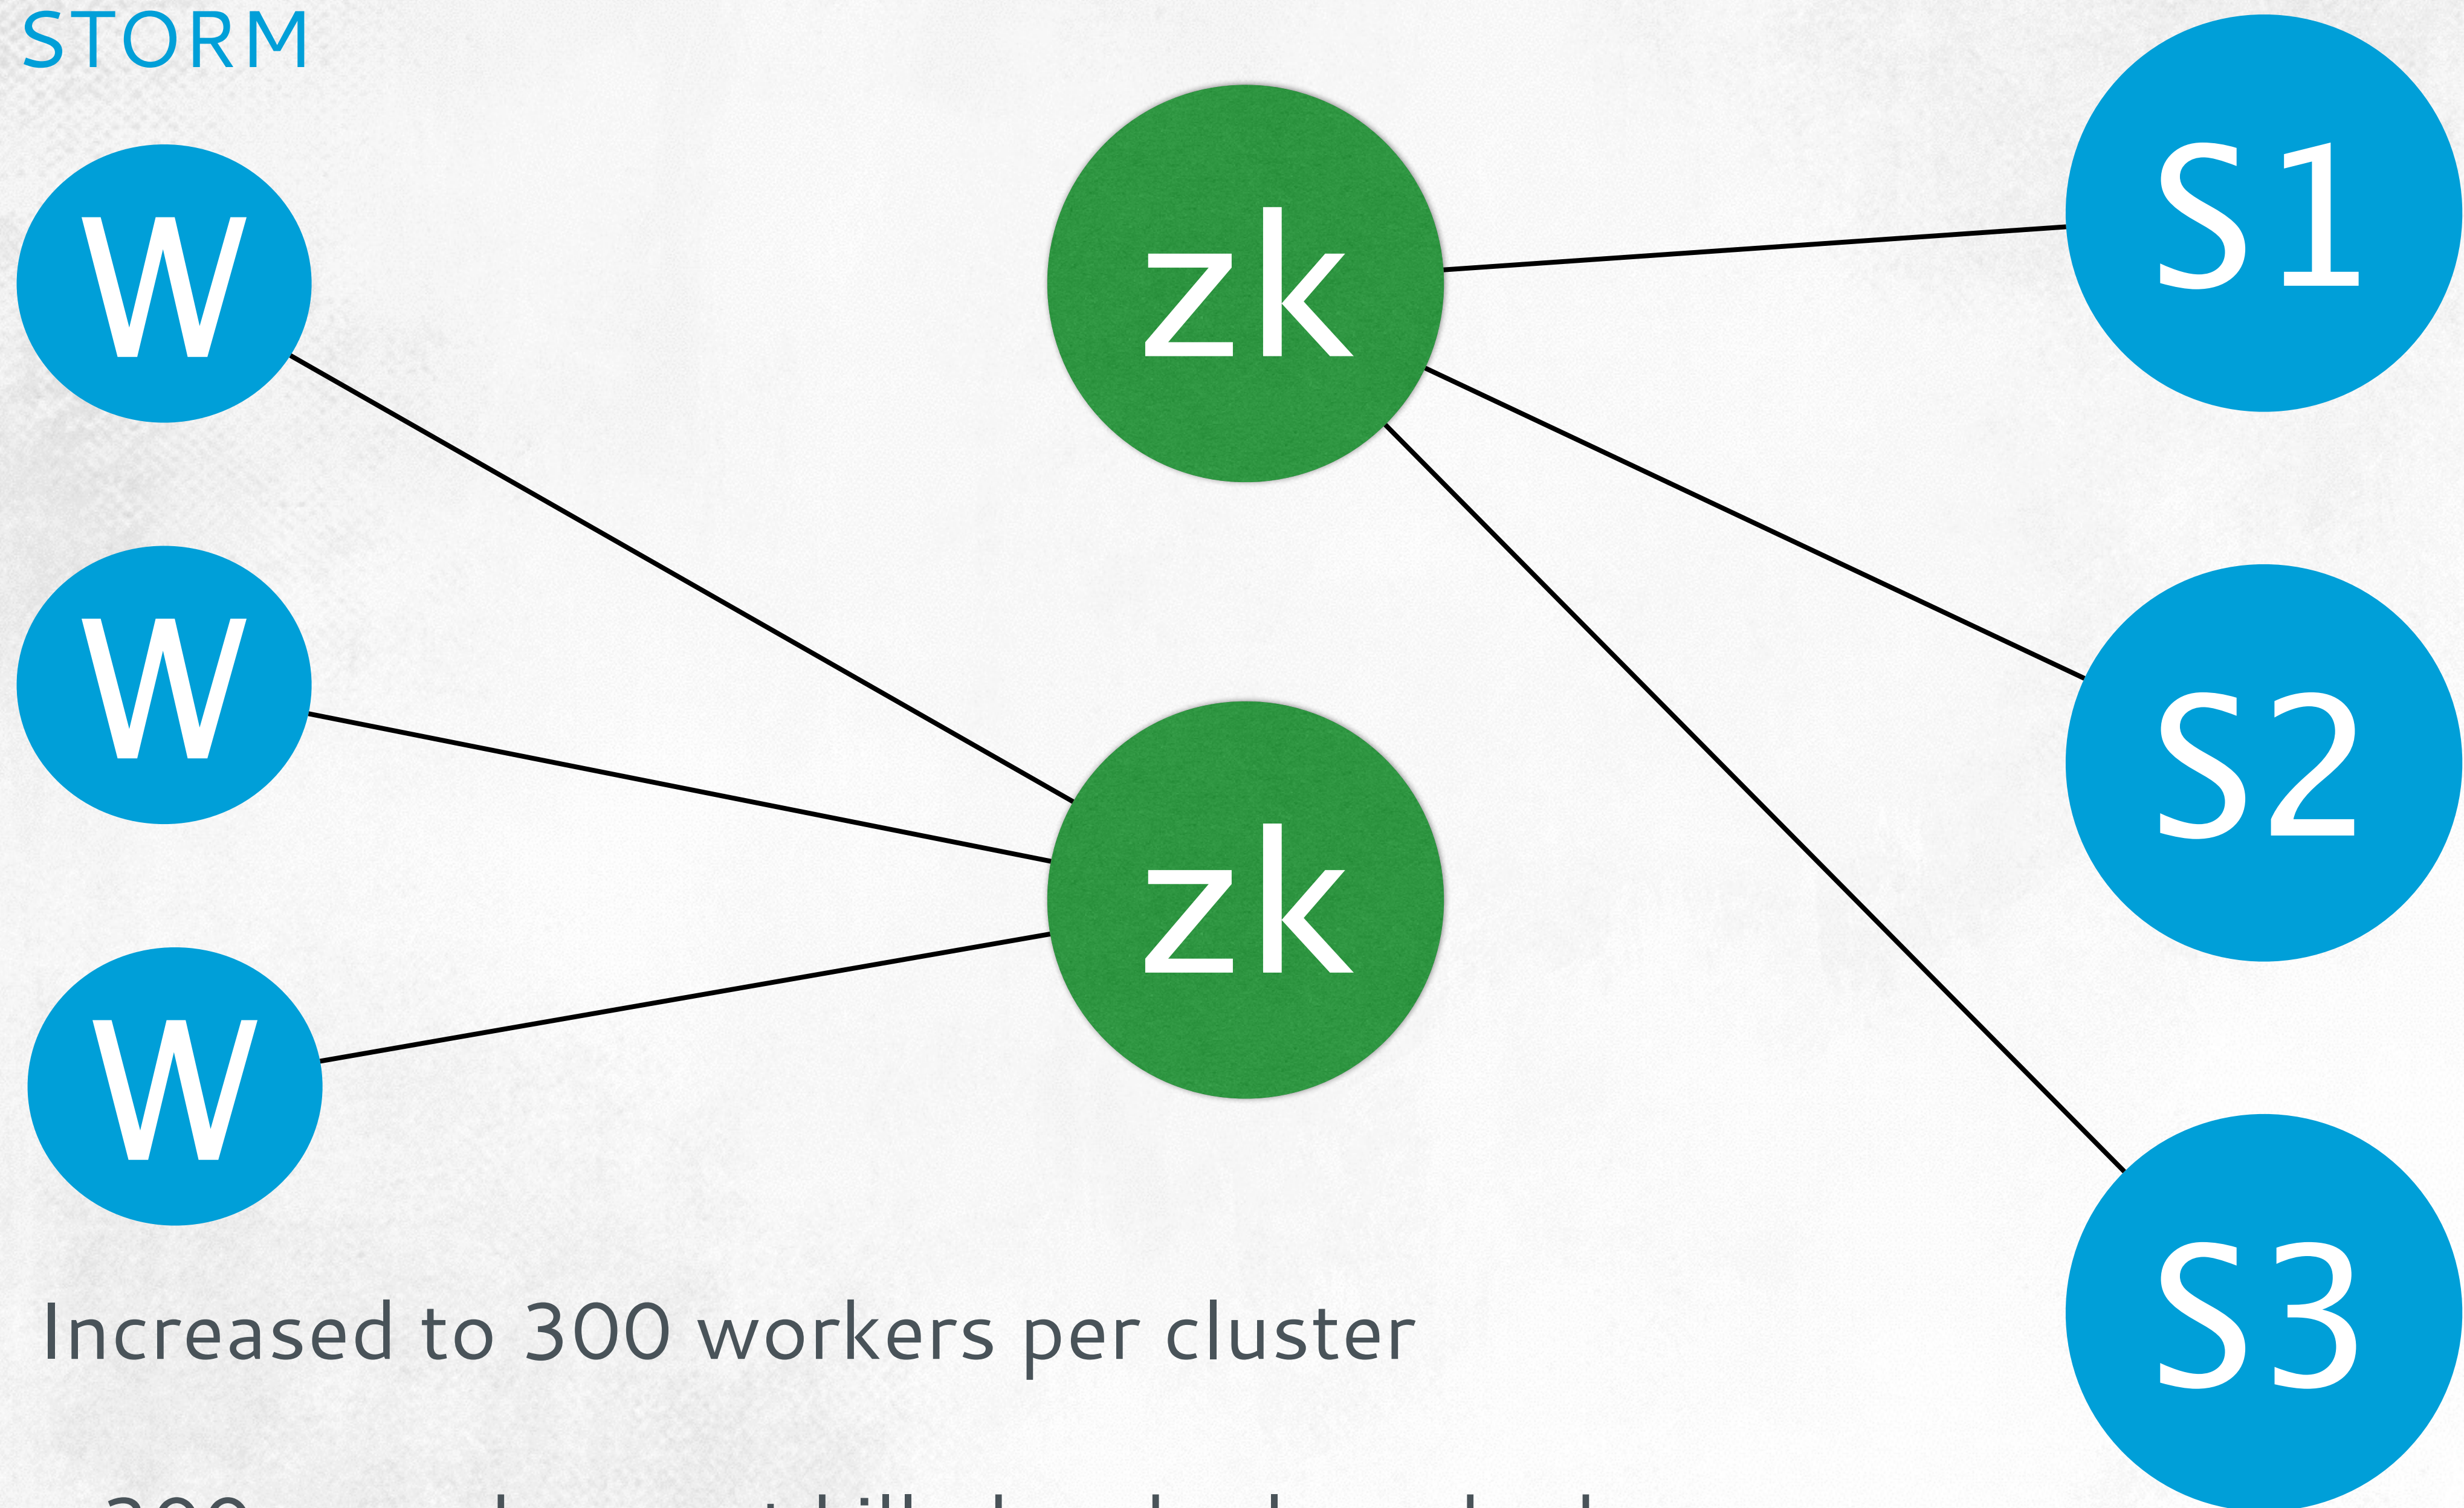
Impacted uptime of other systems



OVERLOADED ZOOKEEPER

Detached configuration

STORM



Increased to 300 workers per cluster

> 300 – workers get killed and relaunched


Worker heart beats written to zknode every 15 secs



OVERLOADED ZOOKEEPER

Scale up

STORM



Increased to 1200 workers per cluster



OVERLOADED ZOOKEEPER

Analyzing zookeeper traffic

67%

KAFKA SPOUT

Offset/partition is written every 2 secs

33%

STORM RUNTIME


Workers write heart beats every 3 secs



OVERLOADED ZOOKEEPER

Heart beat daemons

STORM



5000 workers per cluster and still growing!



STORM OVERHEADS

EXPT 1

JAVA PROGRAM

Read from Kafka cluster and deserialize in a “for loop”

Sustain input rate of 300K msgs/sec from Kafka topic

EXPT 2

1-STAGE TOPOLOGY

No acks to achieve at most once semantics

Storm processes were co-located using isolation scheduler

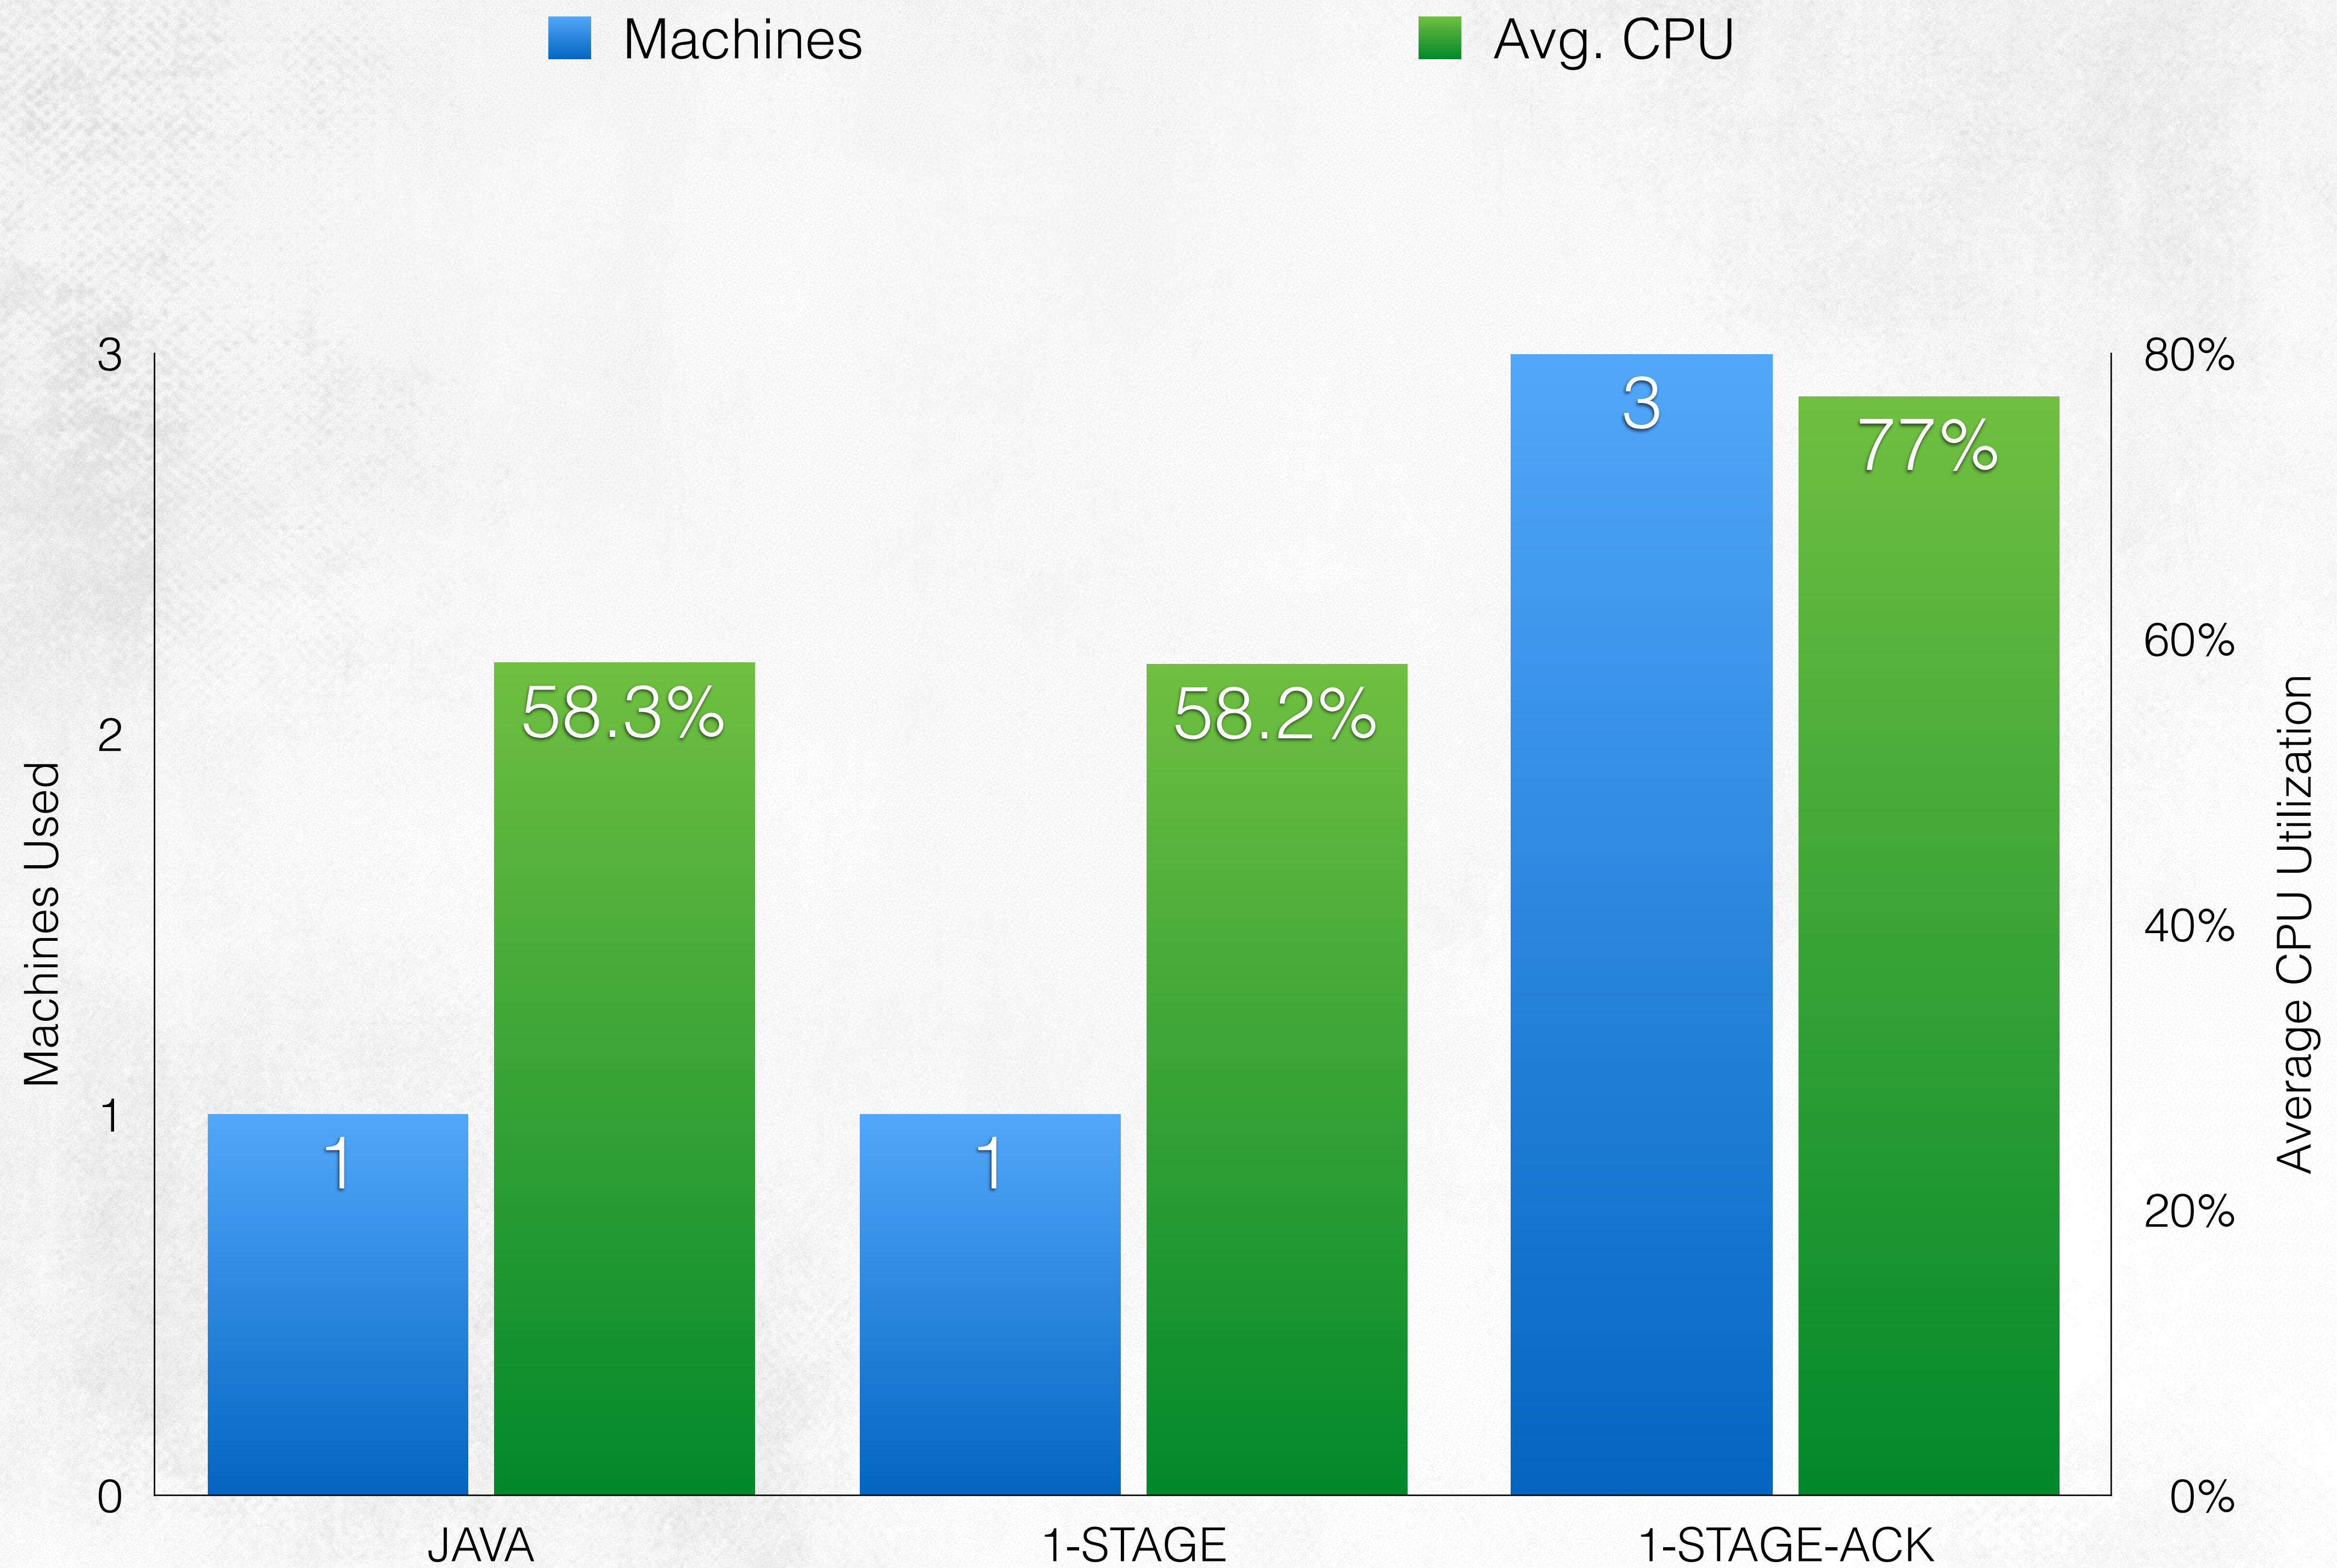
EXPT 3

1-STAGE TOPOLOGY WITH ACKS

Enable acks for at least once semantics



STORM OVERHEADS



STORM EXPERIMENTS

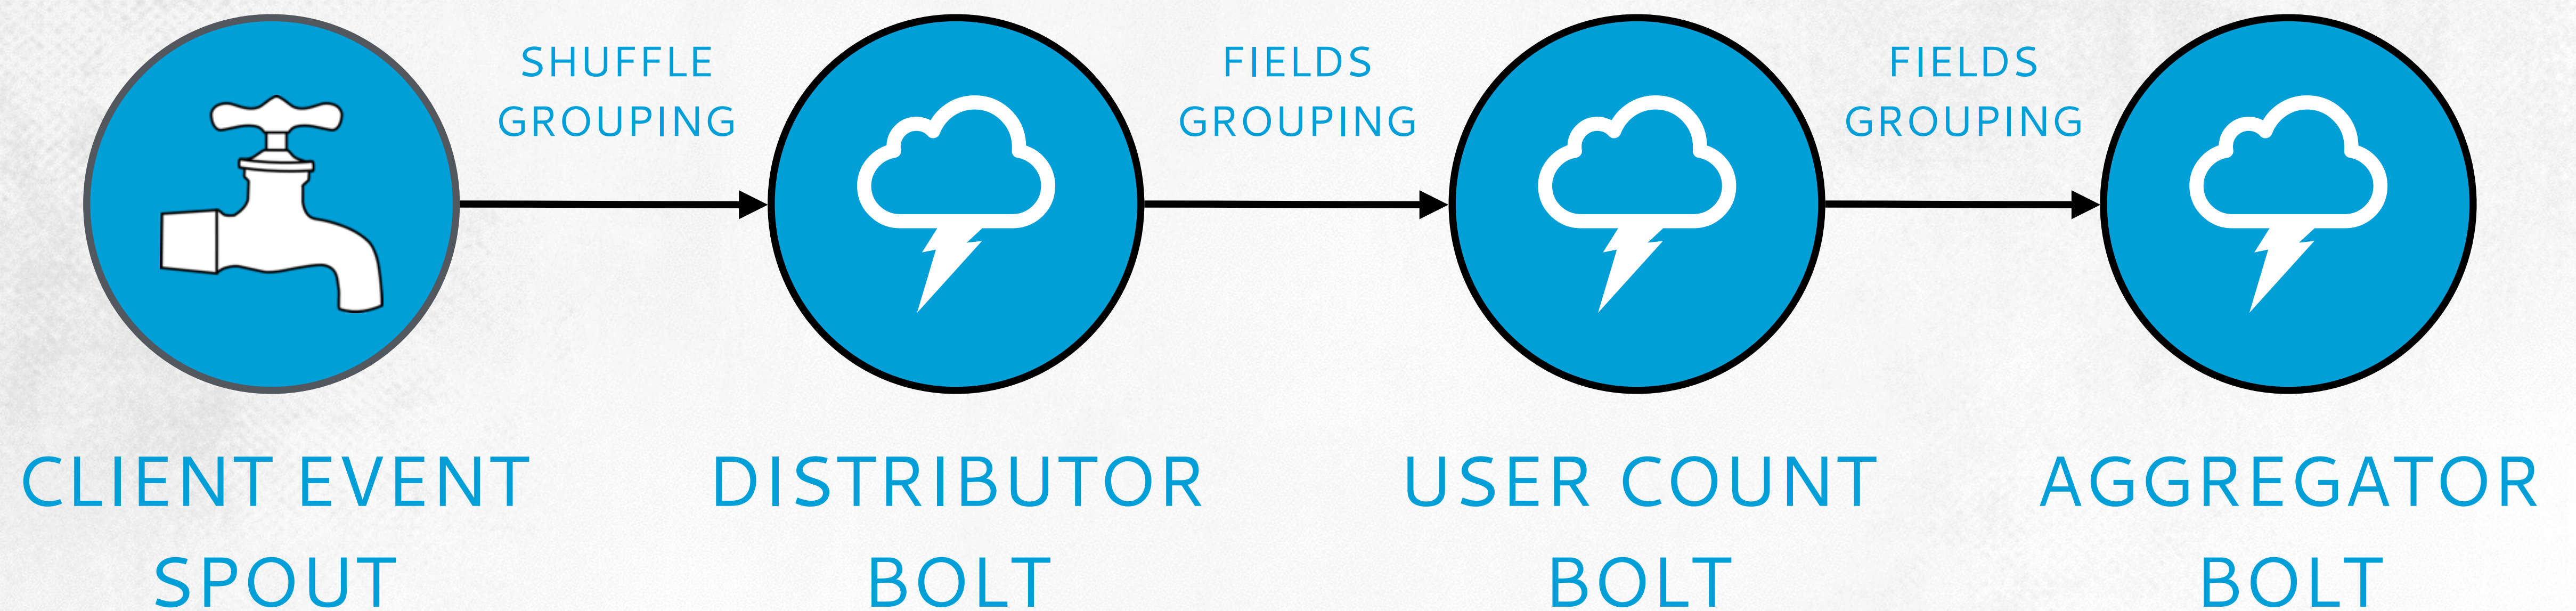


V



STORM EXPERIMENTS


Examine resiliency and efficiency during machine failures





COMPONENTS	# TASKS
client event spout	200
distributor bolt	200
user count bolt	300
aggregator bolt	20



STORM THROUGHPUT



STORM LATENCY



#ThankYou

FOR LISTENING





QUESTIONS AND ANSWERS



Go ahead. Ask away.